

# The Impact of Data Scale and Indexing on Query Performance

---

## Objective

This lab is designed to provide a hands-on understanding of how database performance is affected by two critical factors: **the volume of data** and the use of **database indexes**. You will create a university database, populate it with a growing number of records, measure query execution times, and then analyze the dramatic performance improvements that indexing can provide.

## Tools

- **Database:** PostgreSQL, MySQL, or any other relational database.
- **Programming Language:** Python for data generation and scripting.
- **Python Libraries:**
  - **Faker** : To generate realistic fake data (names, dates, addresses, etc.).
  - **psycopg2-binary** (for PostgreSQL) or **mysql-connector-python** (for MySQL): To connect Python to your database.
  - **matplotlib** or **seaborn** : For creating graphs to visualize your results.

---

## Part 1: Database Schema Design

First, you need to design the structure of your university database. Create a new database (e.g., `university_db`) and then create the following five tables. Pay close attention to the primary keys (PK) and foreign keys (FK) which establish relationships between the tables.

1. **Departments Table** This table stores information about academic departments.

- `department_id` (PK, Integer, Auto-incrementing)
- `department_name` (Varchar) - e.g., 'Computer Science', 'Physics'
- `building` (Varchar) - e.g., 'Block A', 'Science Wing'

2. **Teachers Table** Stores information about the teachers. Each teacher belongs to one department.

- `teacher_id` (PK, Integer, Auto-incrementing)
- `first_name` (Varchar)
- `last_name` (Varchar)
- `email` (Varchar, Unique)
- `department_id` (FK, Integer) - References `Departments.department_id`
- `hire_date` (Date)

3. **Courses Table** Contains details about the courses offered. Each course is taught by a single teacher.

- `course_id` (PK, Integer, Auto-incrementing)
- `course_name` (Varchar) - e.g., 'Introduction to Algorithms', 'Quantum Mechanics'
- `credits` (Integer) - e.g., 3
- `teacher_id` (FK, Integer) - References `Teachers.teacher_id`

4. **Students Table** This is your largest table, storing student information.

- `student_id` (PK, Integer, Auto-incrementing)
- `first_name` (Varchar)
- `last_name` (Varchar)
- `email` (Varchar, Unique)
- `enrollment_date` (Date)
- `date_of_birth` (Date)

5. **Enrollments Table** This is a "junction table" that links students to the courses they are enrolled in, creating a many-to-many relationship.

- `enrollment_id` (PK, Integer, Auto-incrementing)
- `student_id` (FK, Integer) - References `Students.student_id`
- `course_id` (FK, Integer) - References `Courses.course_id`
- `semester` (Varchar) - e.g., 'Fall 2024'
- `grade` (Integer) - A score from 0 to 100.

---

## Part 2: Data Generation and Insertion

Your task is to write a Python script to populate these tables. Use the `Faker` library to generate realistic data.

**Data Generation Plan:**

- **Departments:** 10 departments.
- **Teachers:** 100 teachers, distributed among the 10 departments.
- **Courses:** 200 courses, with each course assigned to a random teacher.
- **Students & Enrollments:** This is where you will scale the data. For each student, enroll them in 5-10 random courses.

Your script should be able to generate data for the following scales:

1. **Scale 1:** 1,000 Students (approx. 5k-10k enrollments)
  2. **Scale 2:** 10,000 Students (approx. 50k-100k enrollments)
  3. **Scale 3:** 100,000 Students (approx. 500k-1M enrollments)
  4. **Scale 4:** 1,000,000 Students (approx. 5M-10M enrollments)
- 

## Part 3: The Queries (To be Timed)

You will run the following five queries at each data scale and record the execution time.

### Query 1: Simple Filter

- **Task:** Find all students who enrolled in the year 2023.
- **Complexity:** A simple WHERE clause on the largest table (`Students`).

### Query 2: Simple Join and Filter

- **Task:** Find the email addresses of all students taught by a specific teacher (e.g., a teacher with `teacher_id = 50`).
- **Complexity:** Involves joining `Students`, `Enrollments`, and `Courses`.

### Query 3: Multi-Join with Text Search

- **Task:** List the full names of all teachers who teach a course with the word 'Advanced' in its name.
- **Complexity:** Involves joining `Teachers` and `Courses` with a LIKE clause for text matching.

### Query 4: Join with Aggregation (GROUP BY)

- **Task:** For each department, count how many courses are being offered. List the department name and the total course count.
- **Complexity:** Requires joining `Departments`, `Teachers`, and `Courses`, and using COUNT with GROUP BY.

### Query 5: Complex Join with Aggregation, Filtering, and Sorting

- **Task:** Find the top 10 students with the highest average grade in the 'Spring 2025' semester. Display the student's full name and their calculated average grade, sorted from highest to lowest.
  - **Complexity:** This is the most challenging query. It requires joining `Students` and `Enrollments`, filtering by semester, grouping by student, calculating an average using AVG, sorting the results, and taking only the top 10 using LIMIT.
- 

## Part 4: The Experiment

Follow these steps methodically.

### Phase 1: Performance Without Indexes

1. Start with a clean, un-indexed set of tables (besides the default Primary Key indexes).
2. **Run for Scale 1 (1k students):**
  - Insert the data.
  - Run each of the 5 queries 3 times and record the average execution time for each.
3. **Run for Scale 2 (10k students):** Repeat the process.
4. **Run for Scale 3 (100k students):** Repeat the process.
5. **Run for Scale 4 (1M students):** Repeat the process.

### Phase 2: Introducing Indexes

Based on the queries you ran, identify the columns that are frequently used in WHERE clauses, JOIN conditions, and ORDER BY clauses. Just for an idea, you might consider indexing the following columns (but feel free to adjust based on your query patterns):

- On the `Students` table: an index on the `enrollment_date` column.
- On the `Enrollments` table: indexes on `student_id`, `course_id`, and `semester`.
- On the `Courses` table: an index on `teacher_id` and `course_name`.
- On the `Teachers` table: an index on `department_id`.

### Phase 3: Performance With Indexes

1. Using your **1 Million student dataset** with the newly created indexes, re-run all 5 queries.
  2. Run each query 3 times and record the average execution time for each.
- 

## Part 5: Analysis and Visualization

Create two graphs to visualize your findings.

### Graph 1: Query Performance vs. Data Scale

- **Type:** Line Chart
- **X-axis:** Data Size (1k, 10k, 100k, 1M)
- **Y-axis:** Execution Time (in milliseconds)

- **Content:** Plot 5 separate lines, one for each query. This graph will visually demonstrate how the execution time for each query grows as the dataset size increases.

#### Graph 2: Impact of Indexing on 1 Million Records

- **Type:** Bar Chart
  - **X-axis:** Query (Query 1, Query 2, Query 3, Query 4, Query 5)
  - **Y-axis:** Execution Time (in milliseconds, consider a logarithmic scale if differences are huge)
  - **Content:** For each query, show two bars side-by-side: one for the time **before indexing** and one for the time **after indexing**. This will powerfully illustrate the performance gain from indexing.
- 

## Part 6: Lab Report

Submit a report that includes the following:

1. Your SQL schema ( `CREATE TABLE` statements).
2. Your Python data generation script.
3. A table containing all the timing data you collected.
4. The two graphs you generated, with clear labels and titles.
5. A conclusion section answering the following questions:
  - Which query was most affected by the increase in data volume? Why do you think that is?
  - Which query saw the most significant performance improvement after indexing? Why?
  - Was there any query that did not improve much with indexing? If so, explain why that might be.
  - What are the potential downsides of adding too many indexes to a database? (Think about `INSERT`, `UPDATE`, and `DELETE` operations).