

LAB: Helpdesk Ticketing System with Redis (Cloud + Insight + Python)

Problem Statement

You are building a Helpdesk Ticketing System for an e-commerce brand. Customers create support tickets (order tracking, returns, technical issues). Agents are assigned based on skills and workload. Each ticket has a priority, status, an SLA due time, and an activity log (events like created, assigned, updated, closed).

You must design this system on **Redis** to achieve real-time performance and simple, scalable data modeling. You will:

1. **Host Redis in the cloud** using Redis Cloud and connect to it via **RedisInsight** (GUI).
2. **Choose correct Redis data structures** for each entity:
 - o **HASH** for structured records (users, agents, tickets),
 - o **SET** for agent skills and per-agent open ticket lists,
 - o **ZSET** for a global priority queue of tickets,
 - o **LIST** for append-only ticket activity logs,
 - o **STRING** for counters and email→user indexes.
3. **Seed a realistic dataset** (users, agents, tickets with mixed priorities/statuses, logs).
4. Perform **CRUD** both manually in RedisInsight and programmatically with Python (**redis-py**).
5. Enforce consistency rules:
 - o If a ticket is assigned to an agent, it must appear in that agent's open_tickets set and increment their load.
 - o Closed tickets must be removed from the global priority queue.
 - o Email lookup must resolve to the correct user.

Your deliverables include screenshots (connection, keyspace, CLI outputs, Python run), and a short write-up explaining why you chose each Redis datatype and how you preserved consistency.

Part A — Setup (Cloud → Insight)

1. **Create a Redis Cloud database**
 - Sign up at redis.com, create a free database.
 - Note the Public Endpoint (host:port), Username (often default), and Password. TLS is enabled by default.
2. **Connect with RedisInsight**
 - Open RedisInsight → Add database → fill:
 - o Host: YOUR_REDIS_HOST

- Port: YOUR_REDIS_PORT
- Username: default (or yours)
- Password: YOUR_REDIS_PASSWORD
- Enable TLS/SSL
- Save & connect. You should see Browser/CLI/Workbench.

Instruction #1 (mandatory): Take a screenshot of the successful RedisInsight connection and include it in your submission

Part B — Scenario & Data Model

Entities & Chosen Data Types

| Entity | Key Pattern | Type | Fields / Members | Rationale |
|-----------------------|------------------------------|---------------|--|--|
| User | help:user:{id} | HASH | name, email, phone, joined_at | Structured attributes; partial updates |
| Email → UserId | help:idx:user_email:{email} | STRING | user_id | Fast lookup by email |
| Agent | help:agent:{id} | HASH | name, email, load | Update individual fields easily |
| Agent Skills | help:agent:{id}:skills | SET | e.g., returns, billing | Membership tests & intersections |
| Ticket | help:ticket:{id} | HASH | user_id, subject, status, priority, created_at, assigned_agent, sla_due_at | Ticket record |
| Ticket Log | help:ticket:{id}:log | LIST | strings like timestamp: action | Append-only event timeline |
| Global Priority Queue | help:queue:priority | ZSET | member=ticket:{id}, score=priority | Sort & pop highest priority |
| Agent's Open Tickets | help:agent:{id}:open_tickets | SET | ticket:{id} | Fast membership/listing |
| ID Counters | help:seq:users | agents | tickets` | STRING (int) |

Instruction #2 (mandatory): Use these data types. If you add more, include a short justification.

Part C — Required Scenarios

CREATE (5 tasks)

1. Create **5 users** + their email index keys.
2. Create **3 agents** + their **skills** sets.
3. Create **8 tickets** with priority mix (1/5/10/20).
4. Assign at least **3 tickets** to agents (update open_tickets + load).
5. Push ≥ 2 log entries per ticket.

READ (6 tasks)

1. Resolve a user by email (index → user_id) and fetch their HASH.
2. Fetch **top-3 highest priority** tickets (ZREVRANGE or ZREVRANGE ... WITHSCORES).
3. List an agent's open tickets (SET) and show each ticket's subject.
4. Display a ticket's log timeline (LIST).
5. Compute counts of tickets by status (client-side tally).
6. Identify agents who cover “**returns**” by inspecting each skills SET.

UPDATE (6 tasks)

1. open → in_progress and log it.
2. in_progress → on_hold and log it.
3. Reassign a ticket A202 → A201 (adjust sets + load).
4. Extend sla_due_at (HSET).
5. Keep load consistent on every assign/unassign.
6. Fix a typo in subject.

DELETE (4 tasks)

1. **Business close:** set status=closed, ZREM from queue, remove from agent open_tickets, log “closed”.
2. **Hard delete** one demo ticket (DEL ticket + log + ZREM).
3. Delete a user and its email index key.
4. Cleanup: use SCAN + pattern to remove lab keys (careful—no KEYS in production).

Instruction #6 (mandatory): For each CRUD category, show all the commands .

Part D — Validation / Integrity Checks

- **Urgent list:**
ZREVRANGE help:queue:priority 0 4 WITHSCORES

- **Agent workload check:**
HGET help:agent:A201 load (compare across agents client-side)
- **Consistency rules:**
 - Assigned ticket must appear in the agent's open_tickets set.
 - Closed tickets must **not** appear in the ZSET queue.
 - Email index must resolve to the correct user.

Part E — Submission Checklist

1. RedisInsight connection screenshot
2. Key Browser snapshot showing your key patterns
3. CLI outputs for CRUD (screenshots)
4. Python run output screenshots
5. Short write-up ($\frac{1}{2}$ –1 page):
 - Why HASH/SET/ZSET/LIST were the right choices
 - Trade-offs (e.g., ZSET vs LIST for priority)
 - How you kept data consistent (load, sets, queue)